

# adanet\_tpu

April 24, 2019

Copyright 2018 The AdaNet Authors.

```
In [0]: #@title Licensed under the Apache License, Version 2.0 (the "License");  
# you may not use this file except in compliance with the License.  
# You may obtain a copy of the License at  
#  
# https://www.apache.org/licenses/LICENSE-2.0  
#  
# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.  
# See the License for the specific language governing permissions and  
# limitations under the License.
```

## 1 AdaNet on TPU

<https://colab.research.google.com/github/tensorflow/adanet/blob/master/...>

<https://github.com/tensorflow/adanet/blob/master/adanet/examples/tutorial/...>

AdaNet supports training on Google's custom machine learning accelerators known as Tensor Processing Units (TPU). Conveniently, we provide `adanet.TPUEstimator` which handles TPU support behind the scenes. There are only a few minor changes needed to switch from `adanet.Estimator` to `adanet.TPUEstimator`. We highlight the necessary changes in this tutorial.

If the reader is not familiar with AdaNet, it is recommended they take a look at [The AdaNet Objective](#) and in particular [Customizing AdaNet](#) as this tutorial builds upon the latter.

**NOTE: you must provide a valid GCS bucket to use `TPUEstimator`.**

To begin, we import the necessary packages, obtain the Colab's TPU master address, and give the TPU permissions to write to our GCS Bucket. Follow the instructions [here](#) to connect to a Colab TPU runtime.

```
In [0]: #@test {"skip": true}  
# If you're running this in Colab, first install the adanet package:  
!pip install adanet  
  
In [0]: from __future__ import absolute_import  
from __future__ import division
```

```

from __future__ import print_function

import functools
import json
import os
import six
import time

import adanet
from google.colab import auth
import tensorflow as tf

BUCKET = '' #@param {type: 'string'}
MODEL_DIR = 'gs://{}/{}'.format(
    BUCKET, time.strftime('adanet-tpu-estimator/%Y-%m-%d-%H-%M-%S'))

MASTER = ''
if 'COLAB_TPU_ADDR' in os.environ:
    auth.authenticate_user()

MASTER = 'grpc://' + os.environ['COLAB_TPU_ADDR']

# Authenticate TPU to use GCS Bucket.
with tf.Session(MASTER) as sess:
    with open('/content/adc.json', 'r') as file_:
        auth_info = json.load(file_)
        tf.contrib.cloud.configure_gcs(sess, credentials=auth_info)

# The random seed to use.
RANDOM_SEED = 42

```

## 1.1 Fashion MNIST

We focus again on the Fashion MNIST dataset and download the data via Keras.

```

In [0]: (x_train, y_train), (x_test, y_test) = (
        tf.keras.datasets.fashion_mnist.load_data())

```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-
32768/29515 [=====] - 0s 0us/step
40960/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-
26427392/26421880 [=====] - 0s 0us/step
26435584/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-
16384/5148 [=====]
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-

```

```
4423680/4422102 [=====] - 0s 0us/step
4431872/4422102 [=====] - 0s 0us/step
```

## 1.2 input\_fn Changes

There are two minor changes we must make to `input_fn` to support running on TPU:

1. TPUs dynamically shard the input data depending on the number of cores used. Because of this, we augment `input_fn` to take a dictionary `params` argument. When running on TPU, `params` contains a `batch_size` field with the appropriate batch size.
2. Once the input is batched, we drop the last batch if it is smaller than `batch_size`. This can simply be done by specifying `drop_remainder=True` to the `tf.data.DataSet.batch()` function. It is important to specify this option since TPUs do not support dynamic shapes. Note that we only drop the remainder batch during training since evaluation is still done on the CPU.

```
In [0]: FEATURES_KEY = "images"
```

```
def generator(images, labels):
    """Returns a generator that returns image-label pairs."""

    def _gen():
        for image, label in zip(images, labels):
            yield image, label

    return _gen

def preprocess_image(image, label):
    """Preprocesses an image for an `Estimator`."""
    image = image / 255.
    image = tf.reshape(image, [28, 28, 1])
    features = {FEATURES_KEY: image}
    return features, label

def input_fn(partition, training, batch_size):
    """Generate an input_fn for the Estimator."""

    def _input_fn(params): # TPU: specify `params` argument.

        # TPU: get the TPU set `batch_size`, if available.
        batch_size_ = params.get("batch_size", batch_size)

        if partition == "train":
            dataset = tf.data.Dataset.from_generator(
```

```

        generator(x_train, y_train), (tf.float32, tf.int32), ((28, 28), ()))
elif partition == "predict":
    dataset = tf.data.Dataset.from_generator(
        generator(x_test[:10], y_test[:10]), (tf.float32, tf.int32),
        ((28, 28), ()))
else:
    dataset = tf.data.Dataset.from_generator(
        generator(x_test, y_test), (tf.float32, tf.int32), ((28, 28), ()))

if training:
    dataset = dataset.shuffle(10 * batch_size_, seed=RANDOM_SEED).repeat()

# TPU: drop the remainder batch when training on TPU.
dataset = dataset.map(preprocess_image).batch(
    batch_size_, drop_remainder=training)
iterator = dataset.make_one_shot_iterator()
features, labels = iterator.get_next()
return features, labels

return _input_fn

```

### 1.3 model\_fn Changes

We use a similar CNN architecture as used in the [Customizing AdaNet](#) tutorial. The only TPU specific change we need to make is wrap the optimizer in a `tf.contrib.tpu.CrossShardOptimizer`.

```

In [0]: #@title Define the Builder and Generator
class SimpleCNNBuilder(adanet.subnetwork.Builder):
    """Builds a CNN subnetwork for AdaNet."""

    def __init__(self, learning_rate, max_iteration_steps, seed):
        """Initializes a `SimpleCNNBuilder`.

        Args:
            learning_rate: The float learning rate to use.
            max_iteration_steps: The number of steps per iteration.
            seed: The random seed.

        Returns:
            An instance of `SimpleCNNBuilder`.
        """

        self._learning_rate = learning_rate
        self._max_iteration_steps = max_iteration_steps
        self._seed = seed

    def build_subnetwork(self,
                        features,
                        logits_dimension,

```

```

        training,
        iteration_step,
        summary,
        previous_ensemble=None):
    """See `adanet.subnetwork.Builder`. """
    images = list(features.values())[0]
    kernel_initializer = tf.keras.initializers.he_normal(seed=self._seed)
    x = tf.keras.layers.Conv2D(
        filters=16,
        kernel_size=3,
        padding="same",
        activation="relu",
        kernel_initializer=kernel_initializer)(
        images)
    x = tf.keras.layers.MaxPool2D(pool_size=2, strides=2)(x)
    x = tf.keras.layers.Flatten()(x)
    x = tf.keras.layers.Dense(
        units=64, activation="relu", kernel_initializer=kernel_initializer)(
        x)

    logits = tf.keras.layers.Dense(
        units=10, activation=None, kernel_initializer=kernel_initializer)(
        x)

    complexity = tf.constant(1)

    return adanet.Subnetwork(
        last_layer=x,
        logits=logits,
        complexity=complexity,
        persisted_tensors={})

def build_subnetwork_train_op(self,
                               subnetwork,
                               loss,
                               var_list,
                               labels,
                               iteration_step,
                               summary,
                               previous_ensemble=None):
    """See `adanet.subnetwork.Builder`. """

    learning_rate = tf.train.cosine_decay(
        learning_rate=self._learning_rate,
        global_step=iteration_step,
        decay_steps=self._max_iteration_steps)
    optimizer = tf.train.MomentumOptimizer(learning_rate, .9)
    # TPU: wrap the optimizer in a CrossShardOptimizer.

```

```

optimizer = tf.contrib.tpu.CrossShardOptimizer(optimizer)
return optimizer.minimize(loss=loss, var_list=var_list)

def build_mixture_weights_train_op(self, loss, var_list, logits, labels,
                                   iteration_step, summary):
    """See `adanet.subnetwork.Builder`. """
    return tf.no_op("mixture_weights_train_op")

@property
def name(self):
    """See `adanet.subnetwork.Builder`. """
    return "simple_cnn"

class SimpleCNNGenerator(adanet.subnetwork.Generator):
    """Generates a `SimpleCNN` at each iteration. """

    def __init__(self, learning_rate, max_iteration_steps, seed=None):
        """Initializes a `Generator` that builds `SimpleCNNs`.

        Args:
            learning_rate: The float learning rate to use.
            max_iteration_steps: The number of steps per iteration.
            seed: The random seed.

        Returns:
            An instance of `Generator`.
        """
        self._seed = seed
        self._dnn_builder_fn = functools.partial(
            SimpleCNNBuilder,
            learning_rate=learning_rate,
            max_iteration_steps=max_iteration_steps)

    def generate_candidates(self, previous_ensemble, iteration_number,
                           previous_ensemble_reports, all_reports):
        """See `adanet.subnetwork.Generator`. """
        seed = self._seed
        # Change the seed according to the iteration so that each subnetwork
        # learns something different.
        if seed is not None:
            seed += iteration_number
        return [self._dnn_builder_fn(seed=seed)]

```

## 1.4 Launch TensorBoard

Let's run [TensorBoard](#) to visualize model training over time. We'll use [ngrok](#) to tunnel traffic to localhost.

The instructions for setting up Tensorboard were obtained from <https://www.dlology.com/blog/quick-guide-to-run-tensorboard-in-google-colab/>

Run the next cells and follow the link to see the TensorBoard in a new tab.

```
In [0]: #@test {"skip": true}

get_ipython().system_raw(
    'tensorboard --logdir {} --host 0.0.0.0 --port 6006 &'
    .format(MODEL_DIR)
)

# Install ngrok binary.
! wget https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.zip
! unzip ngrok-stable-linux-amd64.zip

print("Follow this link to open TensorBoard in a new tab.")
get_ipython().system_raw('./ngrok http 6006 &')
! curl -s http://localhost:4040/api/tunnels | python3 -c "\
    \"import sys, json; print(json.load(sys.stdin)['tunnels'][0]['public_url'])\""
```

## 1.5 Using `adanet.TPUEstimator` to Train and Evaluate

Finally, we switch from `adanet.Estimator` to `adanet.TPUEstimator`. There are two last changes needed:

1. Update the `RunConfig` to be a `tf.contrib.tpu.RunConfig`. We supply the TPU master address and set `iterations_per_loop=200`. This choice is fairly arbitrary in our case. A good practice is to set it to the number of steps in between summary writes and metric evals.
2. Finally, we specify the `use_tpu` and `batch_size` parameters `adanet.TPUEstimator`.

There is an important thing to note about the `batch_size`: each TPU chip consists of 2 cores with 4 shards each. In the [Customizing AdaNet](#) tutorial, a `batch_size` of 64 was used. To be consistent we use the same `batch_size` per shard and drop the number of training steps accordingly. In other words, since we're running on one TPU we set `batch_size=64*8=512` and `train_steps=1000`. In the ideal case, since we drop the `train_steps` by 5x, this means we're **training 5x faster!**

```
In [0]: #@title AdaNet Parameters

LEARNING_RATE = 0.25  #@param {type:"number"}
TRAIN_STEPS = 1000    #@param {type:"integer"}
BATCH_SIZE = 512      #@param {type:"integer"}
ADANET_ITERATIONS = 2  #@param {type:"integer"}

# TPU: switch `tf.estimator.RunConfig` to `tf.contrib.tpu.RunConfig`.
# The main required changes are specifying `tpu_config` and `master`.
config = tf.contrib.tpu.RunConfig(
    tpu_config=tf.contrib.tpu.TPUConfig(iterations_per_loop=200),
    master=MASTER,
```

```

save_checkpoints_steps=200,
save_summary_steps=200,
tf_random_seed=RANDOM_SEED)

head = tf.contrib.estimator.multi_class_head(
    n_classes=10, loss_reduction=tf.losses.Reduction.SUM_OVER_BATCH_SIZE)
max_iteration_steps = TRAIN_STEPS // ADANET_ITERATIONS
# TPU: switch `adanet.Estimator` to `adanet.TPUEstimator`.
try:
    estimator = adanet.TPUEstimator(
        head=head,
        subnetwork_generator=SimpleCNNGenerator(
            learning_rate=LEARNING_RATE,
            max_iteration_steps=max_iteration_steps,
            seed=RANDOM_SEED),
        max_iteration_steps=max_iteration_steps,
        evaluator=adanet.Evaluator(
            input_fn=input_fn("train", training=False, batch_size=BATCH_SIZE),
            steps=None),
        adanet_loss_decay=.99,
        config=config,
        model_dir=MODEL_DIR,
        # TPU: specify `use_tpu` and the batch_size parameters.
        use_tpu=True,
        train_batch_size=BATCH_SIZE,
        eval_batch_size=32)
except tf.errors.InvalidArgumentError as e:
    six.raise_from(
        Exception(
            "Invalid GCS Bucket: you must provide a valid GCS bucket in the "
            "`BUCKET` form field of the first cell."), e)

results, _ = tf.estimator.train_and_evaluate(
    estimator,
    train_spec=tf.estimator.TrainSpec(
        input_fn=input_fn("train", training=True, batch_size=BATCH_SIZE),
        max_steps=TRAIN_STEPS),
    eval_spec=tf.estimator.EvalSpec(
        input_fn=input_fn("test", training=False, batch_size=BATCH_SIZE),
        steps=None,
        start_delay_secs=1,
        throttle_secs=1,
    ))

print("Accuracy:", results["accuracy"])
print("Loss:", results["average_loss"])

```

Accuracy: 0.8913



Loss: 0.298405

## 1.6 Conclusion

That was easy! With very few changes we were able to transform our original estimator into one which can harness the power of TPUs.